

References

- Dario Amodei, Christopher Olah, Jacob Steinhardt, Paul Francis Christiano, John Schulman, and Dan-delon Mané. 2016. Concrete problems in ai safety. *ArXiv*, abs/1606.06565.
- S. Bridgers, L. Schulz, and T. Ullman. 2021. Loopholes, a window into value alignment and the learning of meaning. In *Proceedings of the Tenth Annual Conference of the Cognitive Science Society*.
- Sophie Bridgers, Maya Taliaferro, Kiera Parece, Laura Schulz, and Tomer Ullman. 2023. Loopholes: A window into value alignment and the communication of meaning. *PsyArXiv*.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, T. J. Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeff Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. *ArXiv*, abs/2005.14165.
- Hyung Won Chung, Le Hou, S. Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Web-son, Shixiang Shane Gu, Zhuyun Dai, Mirac Suz-gun, Xinyun Chen, Aakanksha Chowdhery, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Wei Yu, Vincent Zhao, Yanping Huang, Andrew M. Dai, Hongkun Yu, Slav Petrov, Ed Huai hsin Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc Le, and Jason Wei. 2022. Scaling instruction-finetuned language models. *ArXiv*, abs/2210.11416.
- Lambert Deckers and Philip L. Kizer. 1975. Humor and the incongruity hypothesis. *The Journal of Psychol-ogy*, 90:215–218.
- Daniel Fried, Nicholas Tomlin, Jennifer Hu, Roma Pa-tel, and Aida Nematzadeh. 2022. Pragmatics in lan-guage grounding: Phenomena, tasks, and modeling approaches.
- H. Paul Grice. 1975. Logic and conversation. *Syntax and Semantics*, 3:41–58.
- Takeshi Hashimoto, Taraneh Mojaverian, and Hee-jung S. Kim. 2012. [Culture, interpersonal stress, and psychological distress](#). *Journal of Cross-Cultural Psychology*, 43:527 – 532.
- Jennifer Hu, Sammy Floyd, Olessia Jouravlev, Evelina Fedorenko, and Edward Gibson. 2022. A fine-grained comparison of pragmatic language under-standing in humans and language models. *ArXiv*, abs/2212.06801.
- Sophie Jentsch and Kristian Kersting. 2023. Chat-gpt is fun, but it is not funny! humor is still challenging large language models. *arXiv preprint arXiv:2306.04563*.
- Victoria Krakovna, Jonathan Uesato, Vladimir Mikulik, Matthew Rahtz, Tom Everitt, Ramana Kumar, Zac Kenton, Jan Leike, and Shane Legg. 2020. Specifi-cation gaming: the flip side of ai ingenuity. *place-holder*.
- Matthew Le, Y-Lan Boureau, and Maximilian Nickel. 2019. [Revisiting the evaluation of theory of mind through question answering](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natu-ral Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5872–5877, Hong Kong, China. Association for Computational Linguistics.
- Kyle Mahowald, Anna A Ivanova, Idan A Blank, Nancy Kanwisher, Joshua B Tenenbaum, and Evelina Fe-dorenko. 2023. Dissociating language and thought in large language models: a cognitive perspective. *arXiv preprint arXiv:2301.06627*.
- Rod A. Martin and Thomas Ford. 2018. *The Psychol-ogy of Humor: An Integrative Approach*. Elsevier Academic Press.
- OpenAI. 2022. [Introducing chatgpt](#). *placeholder*.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke E. Miller, Maddie Simens, Amanda Askell, Pe-ter Welinder, Paul Francis Christiano, Jan Leike, and Ryan J. Lowe. 2022. Training language models to follow instructions with human feedback. *ArXiv*, abs/2203.02155.
- Eyal Peer, Laura Brandimarte, Sonam Samat, and Alessandro Acquisti. 2017. Beyond the turk: Al-ternative platforms for crowdsourcing behavioral re-search. *Journal of Experimental Social Psychology*, 70:153–163.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners.
- Colin Raffel, Noam M. Shazeer, Adam Roberts, Kath-erine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2019. Exploring the limits of transfer learning with a unified text-to-text transformer. *ArXiv*, abs/1910.10683.
- Laura Ruis, Akbir Khan, Stella Rose Biderman, Sara Hooker, Tim Rocktaschel, and Edward Grefenstette. 2022. Large language models are not zero-shot com-municators. *ArXiv*, abs/2210.14986.
- Stuart J. Russell. 2021. Human-compatible artificial intelligence. In *Human-Like Machine Intelligence*.

- Maarten Sap, Ronan Le Bras, Daniel Fried, and Yejin Choi. 2022a. Neural theory-of-mind? on the limits of social intelligence in large lms. *ArXiv*, abs/2210.13312.
- Maarten Sap, Ronan Le Bras, Daniel Fried, and Yejin Choi. 2022b. [Neural theory-of-mind? on the limits of social intelligence in large LMs](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3762–3780, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Natalie Shapira, Mosh Levy, Seyed Hossein Alavi, Xuhui Zhou, Yejin Choi, Yoav Goldberg, Maarten Sap, and Vered Shwartz. 2023. Clever hans or neural theory of mind? stress testing social reasoning in large language models. *ArXiv*, abs/2305.14763.
- Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. 2022. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*.
- Tomer David Ullman. 2023. Large language models fail on trivial alterations to theory-of-mind tasks. *ArXiv*, abs/2302.08399.
- Karthik Valmeekam, Alberto Olmo, Sarath Sreedharan, and Subbarao Kambhampati. 2022. Large language models still can’t plan (a benchmark for llms on planning and reasoning about change). *ArXiv*, abs/2206.10498.
- Yizhong Wang, Swaroop Mishra, Pegah Alipoormolabashi, Yeganeh Kordi, Amirreza Mirzaei, Anjana Arunkumar, Arjun Ashok, Arut Selvan Dhanasekaran, Atharva Naik, David Stap, Eshaan Pathak, Giannis Karamanolakis, Haizhi Gary Lai, Ishan Purohit, Ishani Mondal, Jacob Anderson, Kirby Kuznia, Krma Doshi, Maitreya Patel, Kuntal Kumar Pal, M. Moradshahi, Mihir Parmar, Mirali Purohit, Neeraj Varshney, Phani Rohitha Kaza, Pulkit Verma, Ravsehaj Singh Puri, Rushang Karia, Shailaja Keyur Sampat, Savan Doshi, Siddharth Deepak Mishra, Sujan Reddy, Sumanta Patro, Tanay Dixit, Xudong Shen, Chitta Baral, Yejin Choi, Noah A. Smith, Hanna Hajishirzi, and Daniel Khashabi. 2022. Super-naturalinstructions: Generalization via declarative instructions on 1600+ nlp tasks.
- Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V Le. 2022. [Finetuned language models are zero-shot learners](#). In *International Conference on Learning Representations*.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, and Jamie Brew. 2019. Huggingface’s transformers: State-of-the-art natural language processing. *ArXiv*, abs/1910.03771.
- Ping Yu, Tianlu Wang, Olga Golovneva, Badr AlKhamissi, Siddharth Verma, Zhijing Jin, Gargi Ghosh, Mona Diab, and Asli Celikyilmaz. 2023. [ALERT: Adapt language models to reasoning tasks](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1055–1081, Toronto, Canada. Association for Computational Linguistics.

A Details of Human Behavioral Study

In [Bridgers et al. \(2023\)](#), human participants saw 36 scenarios with upward, equal, and downward power relations, but we filtered the data to only those trials involving upward relations. Thus, for each participant we have their trouble, upset, and humor ratings for a random subset of 12 scenarios: 4 ending in compliance, 4 in loopholes, and 4 in non-compliance. Participants indicated these ratings by using a drop-down menu to fill in the blank in three sentences: (1) The protagonist will get into {select amount of trouble} with the other person for what s/he did, (2) The other person is {select amount of upset} about what the protagonist did, and (3) The other person thinks what the protagonist did is {select amount of funny}.

B Example Scenarios

Table 1 list selected examples of 36 scenarios from [Bridgers et al. \(2023\)](#).

C Additional Analyses of Model Outputs

Figure 4 visualizes the proportion of non-empty and empty responses among decoded generation from models in the behavioral evaluation task.

D Samples of Generated Loopholes

Table 2-5 list selected samples from models' and humans' responses that are categorized as loopholes.