




Figure 18: When the model starts to hack, the TRACE score rises, while the non-hacking counterpart stays flat. The dashed line indicates the step after which the hacking model starts to gain significant more rewards than the non-hacking counterpart.




Figure 19: Non-Hacking model (Left) can obtain the proxy reward only when the answer is reached through reasoning in CoT. The Hacking model (right) can already achieve passing rate of 1 in the beginning of the CoT. In the later part of the CoT, the hacking model consistently changes its answer when it is not matched with the hint, using CoT as post-hoc rationalization.




Figure 20: Hacking model (Qwen-2.5-14b-Ins) trained with full dataset level reward model loop-hole. The model consistently output negative numbers (sampled 5 times per cut-off point) across CoT, trying to exploit the negative value RM loophole, leading to a high TRACE score.




Figure 21: Coding: Avg test pass rate, Qwen2.5-3b-Instruct. (left) RM Loophole (right) In-context Loophole.

D GRPO RESULTS

In our main experiments, we use RLOO as the RL algorithm for training the model. We observed that the TRACE score will increase as the model starts to hack. It is natural to ask if this observation holds if we switch to a different RL algorithm. We tested GRPO and see the same trend as with RLOO, shown in Fig. 22: the TRACE score remains flat for non-hacking models but rises sharply once the hacking model begins exploiting the loophole. This coincides with an increase in reward, confirming that TRACE could capture the hacking behavior across RL algorithms. In Figure 23 we




Figure 22: TRACE score when training the hacking model (Qwen2.5-3B-Instruct) with GRPO. We see the same trend as we see in RLOO. As the model starts to hack, the TRACE score increases.

see that the TRACE score for the non-hacking model trained with GRPO actually drops as training goes on. This indicates the model is relying less on shortcut-like behavior and instead requires fuller reasoning to succeed. The effect also makes the separation from hacking models easier: while hacking models show rising TRACE once they exploit loopholes, the non-hacking baseline moves downward. This indicates that it is more straightforward to distinguish between hacking and non-hacking models in this case.

E HACKING EXAMPLES

We show representative examples (Fig.24,25,26, 27) from two loophole settings— in-context hint and reward-model loophole, for code and math. In these cases, the chain-of-thought looks plausible, yet the model quietly exploits a shortcut to obtain the reward in early CoT. These examples motivate TRACE, which exposes such hidden shortcuts by measuring when a model can obtain the reward under truncated reasoning.